



Entregable 4

Silvia Julià

Proyecto I

```
model = ks.Sequential()  
  
model.add(ks.layers.Conv2D(32, (3, 3), strides=1, activation='relu',  
                         padding='same', input_shape=(32,32,3)))  
model.add(ks.layers.MaxPooling2D((2, 2)))  
  
model.add(ks.layers.Flatten())  
model.add(ks.layers.Dense(500, activation='relu'))  
model.add(ks.layers.Dropout(0.5))  
model.add(ks.layers.Dense(10, activation='softmax'))
```

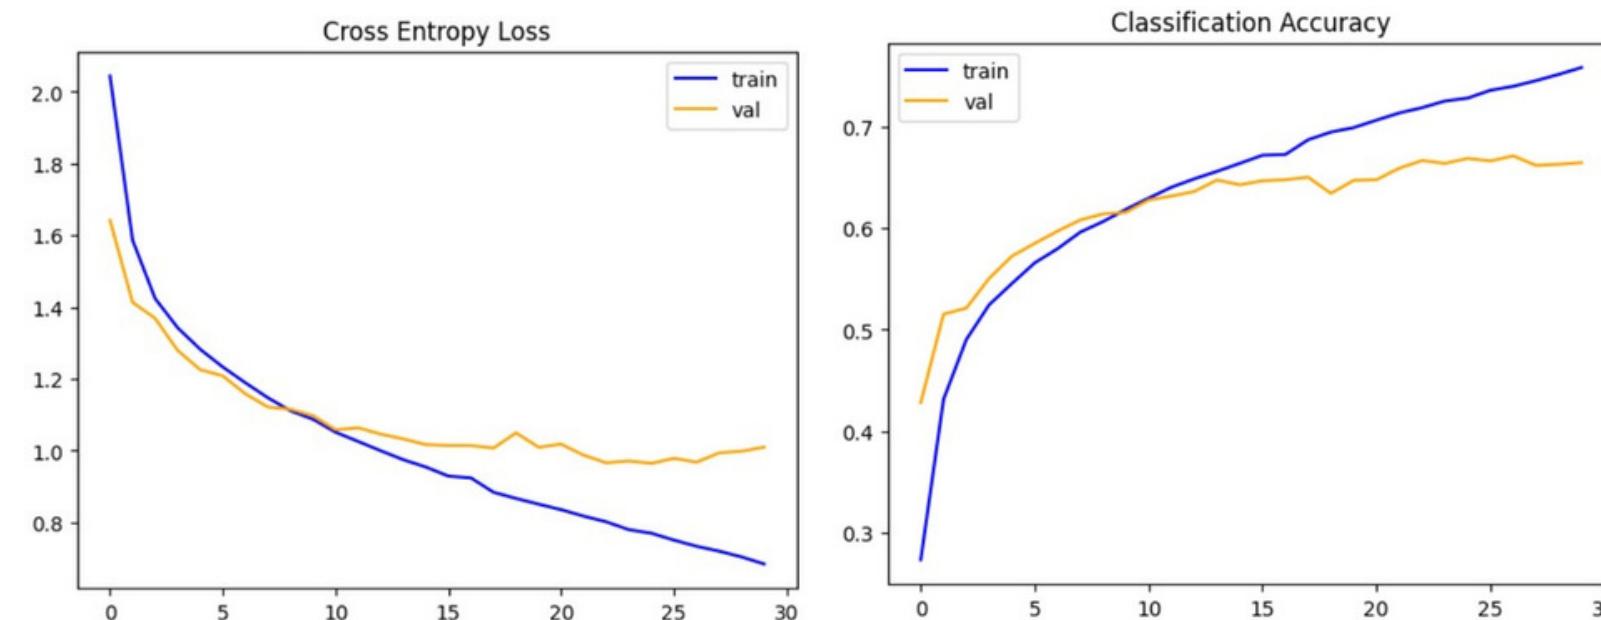
Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 500)	4096500
dropout (Dropout)	(None, 500)	0
dense_1 (Dense)	(None, 10)	5010
<hr/>		
Total params: 4102406 (15.65 MB)		
Trainable params: 4102406 (15.65 MB)		
Non-trainable params: 0 (0.00 Byte)		

```
history = model.fit(x_train_scaled, y_train,  
                     callbacks=[modelchk, callback_loss, callback_acc],  
                     epochs=300, batch_size= 512,  
                     validation_data=(x_val_scaled, y_val))
```

```
_ , acc = model.evaluate(x_test_scaled, y_test, verbose=0)  
print('> %.3f' % (acc * 100.0))
```

> 65.360



Conclusiones y observaciones:

En este primer proyecto aumentamos neuronas en la capa dense, y añadimos un dropout para controlar un poco el overfitting que puede causar este aumento.

El resultado es un accuracy mucho mayor (el proyecto base tenía 60.260) aunque aun falta mejorar tanto el accuracy como el overfitting.

Proyecto II

```
model = Sequential()

# Primera capa de convolución
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same', input_shape=(32,32,3)))
model.add(MaxPooling2D((2, 2)))

# Segunda capa de convolución
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))

# Tercera capa de convolución
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))

model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dropout(0.5))

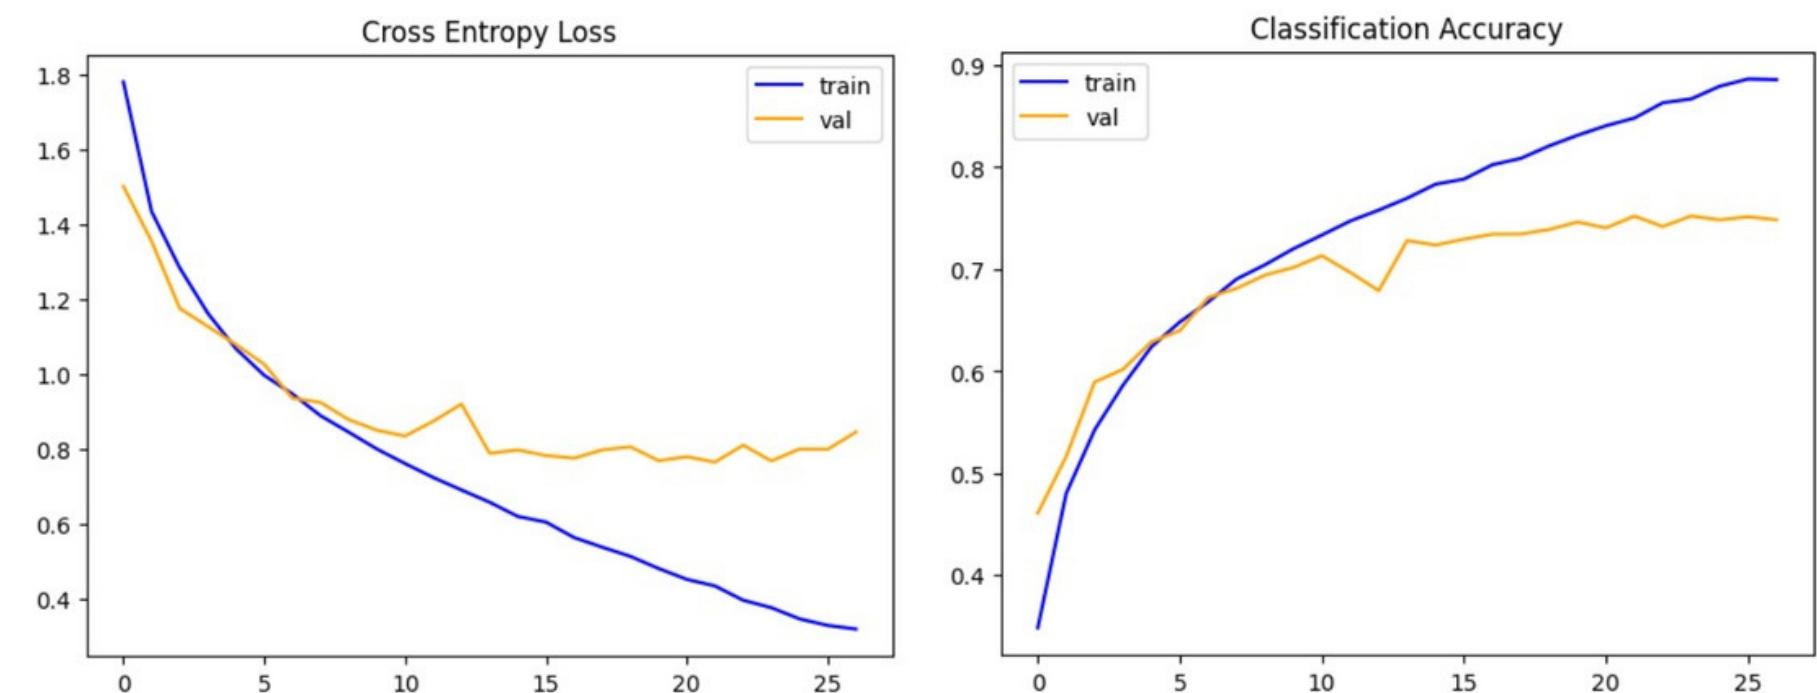
model.add(Dense(10, activation='softmax'))

Layer (type)          Output Shape         Param #
=====
conv2d_10 (Conv2D)    (None, 32, 32, 32)   896
max_pooling2d_10 (MaxPooling2D) (None, 16, 16, 32) 0
conv2d_11 (Conv2D)    (None, 16, 16, 64)    18496
max_pooling2d_11 (MaxPooling2D) (None, 8, 8, 64) 0
conv2d_12 (Conv2D)    (None, 8, 8, 128)    73856
max_pooling2d_12 (MaxPooling2D) (None, 4, 4, 128) 0
flatten_8 (Flatten)   (None, 2048)        0
dense_30 (Dense)     (None, 500)         1024500
dropout_22 (Dropout) (None, 500)         0
dense_31 (Dense)     (None, 10)          5010
=====
Total params: 1122758 (4.28 MB)
Trainable params: 1122758 (4.28 MB)
Non-trainable params: 0 (0.00 Byte)
```

```
history = model.fit(x_train_scaled, y_train,
                     callbacks=[modelchk, callback_loss, callback_acc],
                     epochs=300, batch_size= 512,
                     validation_data=(x_val_scaled, y_val))

_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 73.960



Conclusiones y observaciones:

Con la intención de mejorar el modelo añadimos 2 capas de convulsiones, lo que conlleva una mejora innegable del accuracy, pero un mayor problema de overfitting.

Proyecto III

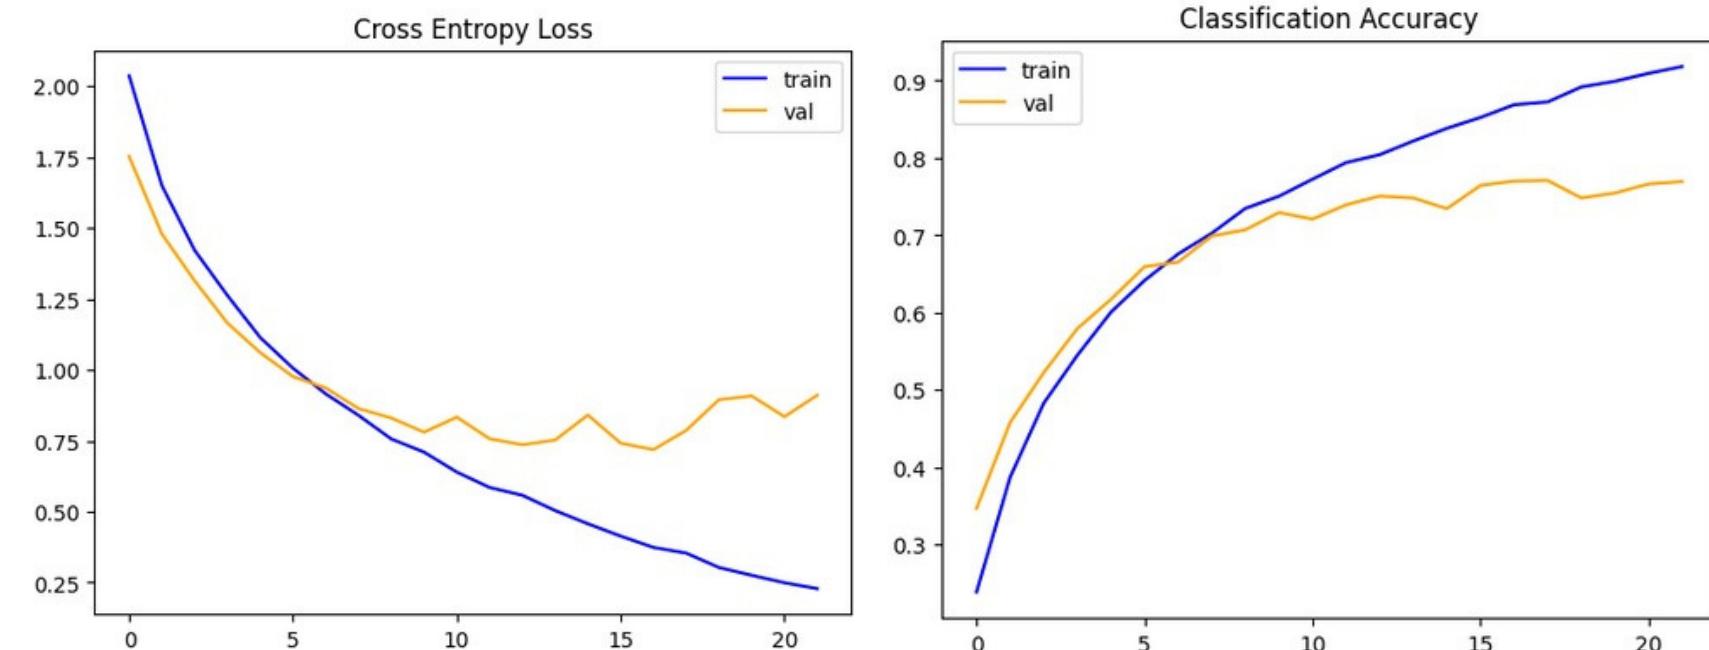
```
model = Sequential()
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same', input_shape=(32,32,3)))
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))
```

=====					
conv2d (Conv2D)	(None, 32, 32, 32)	896	conv2d_6 (Conv2D)	(None, 8, 8, 128)	73856
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248	conv2d_7 (Conv2D)	(None, 8, 8, 128)	147584
conv2d_2 (Conv2D)	(None, 32, 32, 32)	9248	conv2d_8 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0	max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	18496	flatten (Flatten)	(None, 2048)	0
conv2d_4 (Conv2D)	(None, 16, 16, 64)	36928	dense (Dense)	(None, 500)	1024500
conv2d_5 (Conv2D)	(None, 16, 16, 64)	36928	dropout (Dropout)	(None, 500)	0
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0	dense_1 (Dense)	(None, 10)	5010
=====					
			Total params:	1510278 (5.76 MB)	
			Trainable params:	1510278 (5.76 MB)	

```
history = model.fit(x_train_scaled, y_train,
                     callbacks=[modelchk, callback_loss, callback_acc],
                     epochs=300, batch_size= 512,
                     validation_data=(x_val_scaled, y_val))
```

```
, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 76.470



Conclusiones y observaciones:

En este caso decidí probar aumentar las convulsiones d para ver como afecta al modelo.

Aunque como podemos ver seguimos teniendo problemas de overfitting que intentaremos mejorar en el siguiente proyecto.

Proyecto IV

```
model = Sequential()
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same', input_shape=(32,32,3)))
model.add(Dropout(0.4))
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Dropout(0.4))
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))

model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Dropout(0.4))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Dropout(0.4))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))

model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Dropout(0.4))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Dropout(0.4))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))

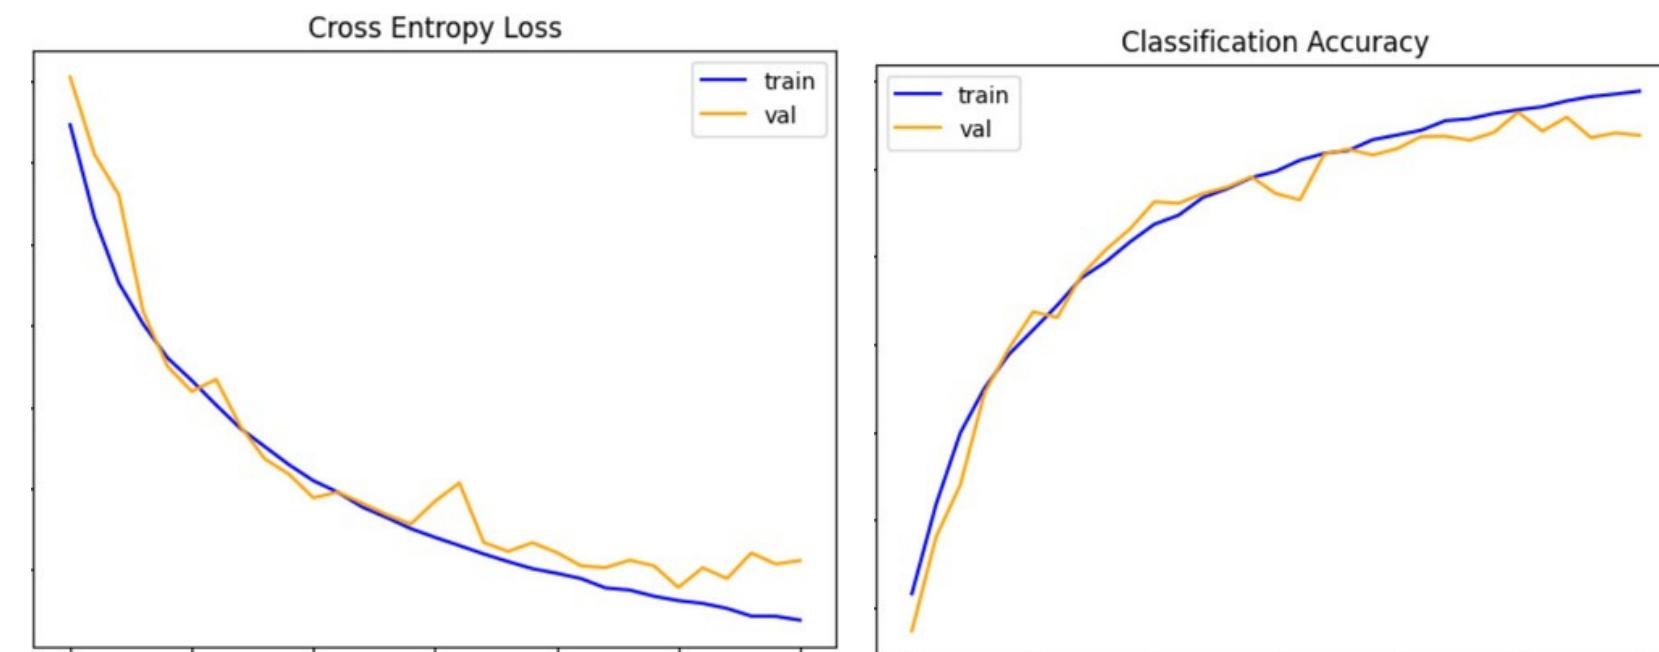
model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))
```


	conv2d_9 (Conv2D)	(None, 32, 32, 32)	896	conv2d_15 (Conv2D)	(None, 8, 8, 128)	73856
dropout_10 (Dropout)	(None, 32, 32, 32)	0	dropout_14 (Dropout)	(None, 8, 8, 128)	0	
conv2d_10 (Conv2D)	(None, 32, 32, 32)	9248	conv2d_16 (Conv2D)	(None, 8, 8, 128)	147584	
dropout_11 (Dropout)	(None, 32, 32, 32)	0	dropout_15 (Dropout)	(None, 8, 8, 128)	0	
conv2d_11 (Conv2D)	(None, 32, 32, 32)	9248	conv2d_17 (Conv2D)	(None, 8, 8, 128)	147584	
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 32)	0	max_pooling2d_5 (MaxPooling2D)	(None, 4, 4, 128)	0	
conv2d_12 (Conv2D)	(None, 16, 16, 64)	18496	flatten_1 (Flatten)	(None, 2048)	0	
dropout_12 (Dropout)	(None, 16, 16, 64)	0	dense_2 (Dense)	(None, 500)	1024500	
conv2d_13 (Conv2D)	(None, 16, 16, 64)	36928	dropout_16 (Dropout)	(None, 500)	0	
dropout_13 (Dropout)	(None, 16, 16, 64)	0	dense_3 (Dense)	(None, 10)	5010	
conv2d_14 (Conv2D)	(None, 16, 16, 64)	36928				
max_pooling2d_4 (MaxPooling2D)	(None, 8, 8, 64)	0	Total params:	1510278 (5.76 MB)		
			Trainable params:	1510278 (5.76 MB)		

```
history = model.fit(x_train_scaled, y_train,
                     callbacks=[modelchk, callback_loss, callback_acc],
                     epochs=300, batch_size= 512,
                     validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 73.480



Conclusiones y observaciones:

Al añadir dropouts de 0.4 parece que ha mejorado el overfitting, pero ha dificultado demasiado el aprendizaje del modelo, pasando de un 76 a un 73. Veamos si se puede mejorar.

Proyecto V

```
model = Sequential()
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same', input_shape=(32,32,3)))
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

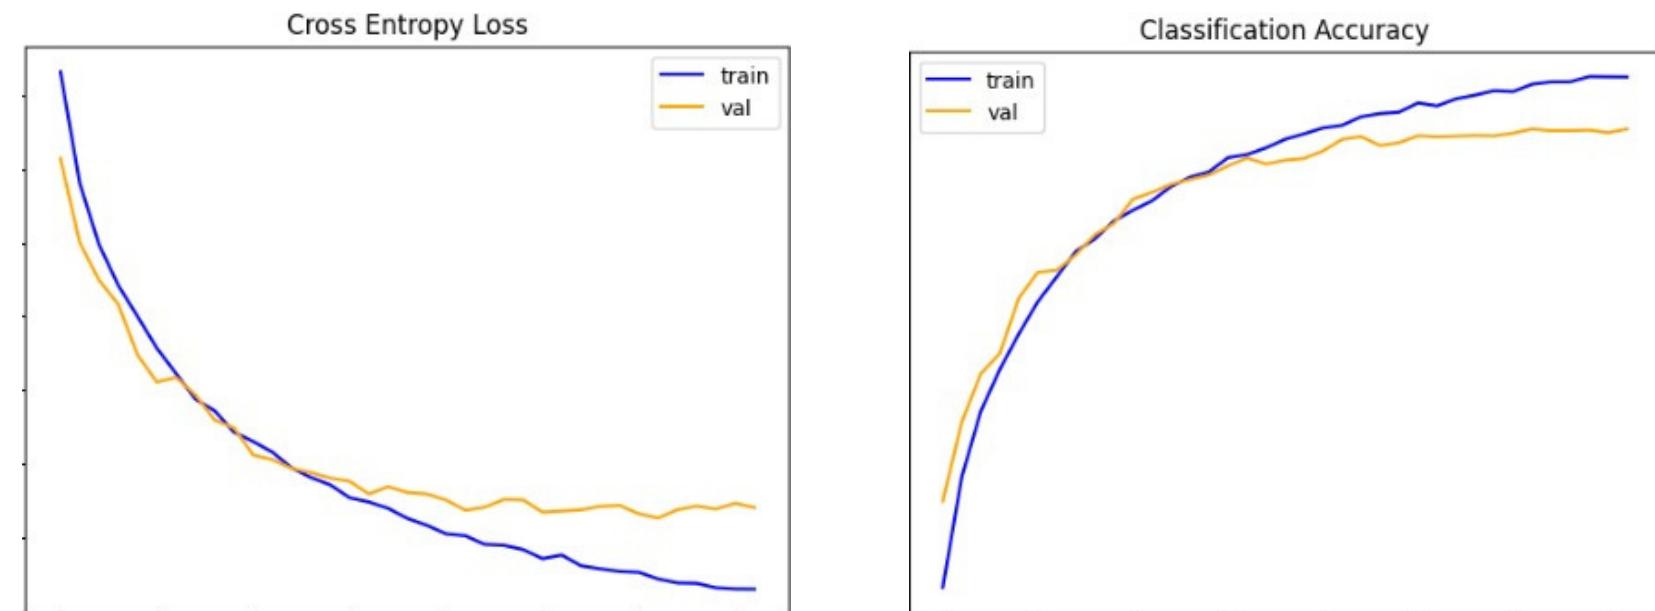
model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))
```

```
=====
conv2d (Conv2D)           (None, 32, 32, 32)    896      conv2d_6 (Conv2D)       (None, 8, 8, 128)    73856
conv2d_1 (Conv2D)          (None, 32, 32, 32)    9248     conv2d_7 (Conv2D)       (None, 8, 8, 128)    147584
conv2d_2 (Conv2D)          (None, 32, 32, 32)    9248     conv2d_8 (Conv2D)       (None, 8, 8, 128)    147584
max_pooling2d (MaxPooling2D) (None, 16, 16, 32)    0        max_pooling2d_2 (MaxPooling2D) (None, 4, 4, 128)    0
dropout (Dropout)          (None, 16, 16, 32)    0        dropout_2 (Dropout)     (None, 4, 4, 128)    0
conv2d_3 (Conv2D)          (None, 16, 16, 64)   18496    flatten (Flatten)     (None, 2048)      0
conv2d_4 (Conv2D)          (None, 16, 16, 64)   36928    dense (Dense)        (None, 500)      1024500
conv2d_5 (Conv2D)          (None, 16, 16, 64)   36928    dropout_3 (Dropout)   (None, 500)      0
max_pooling2d_1 (MaxPooling2D) (None, 8, 8, 64)   0        dense_1 (Dense)       (None, 10)      5010
dropout_1 (Dropout)        (None, 8, 8, 64)      0
=====
Total params: 1510278 (5.76 MB)
Trainable params: 1510278 (5.76 MB)
```

```
history = model.fit(x_train_scaled, y_train,
                     callbacks=[modelchk, callback_loss, callback_acc],
                     epochs=300, batch_size= 512,
                     validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 80.380



Conclusiones y observaciones:

En este caso he decidido dejar a 0.3 los dropouts y al mismo tiempo disminuir la cantidad de estos, ha resultado mejorar tanto el overfitting como el accuracy. Hemos llegado por primera vez a 80.

Proyecto VI

```
model = Sequential()
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same', input_shape=(32, 32, 3))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.25))

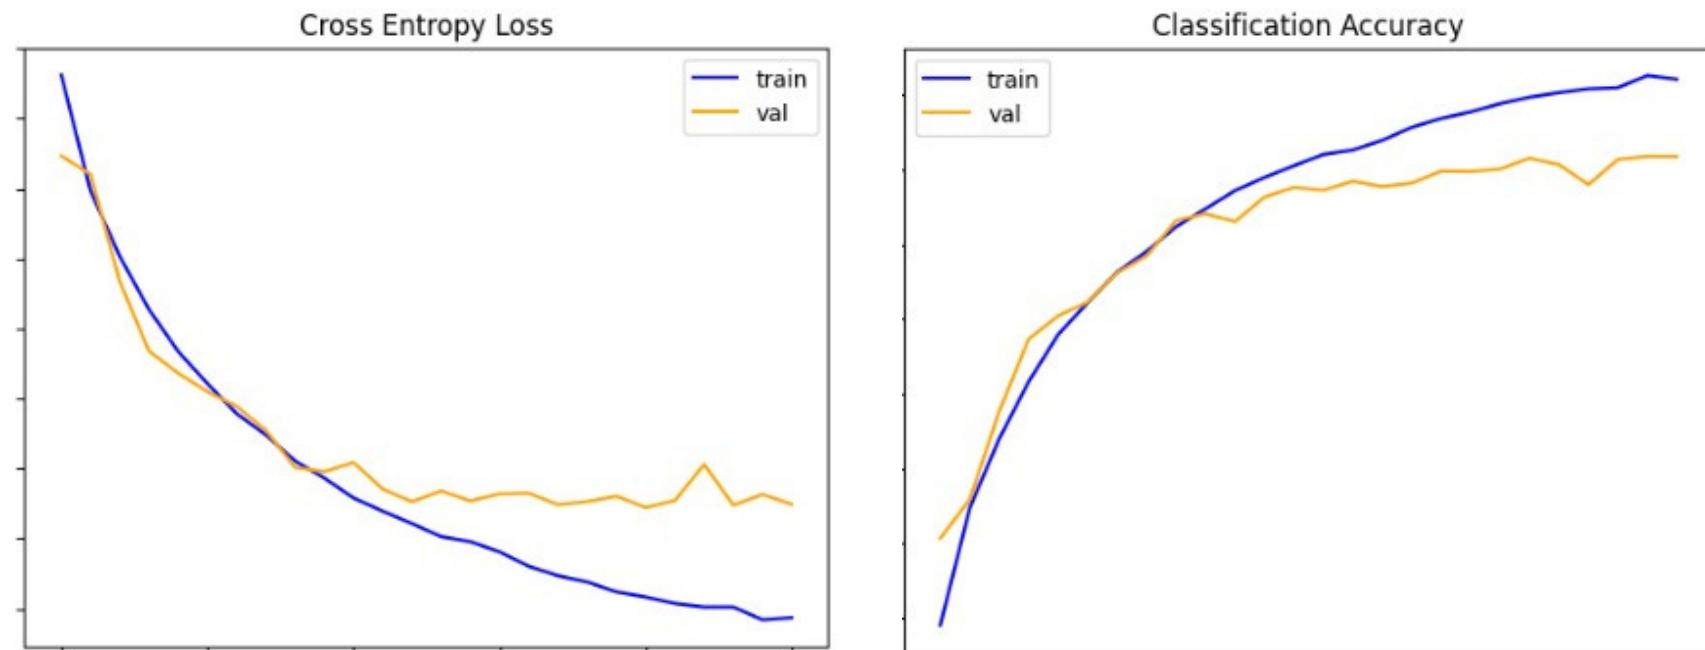
model.add(Conv2D(256, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(256, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(256, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(1024, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))
```

```
history = model.fit(x_train_scaled, y_train,
                     callbacks=[modelchk, callback_loss, callback_acc],
                     epochs=300, batch_size= 512,
                     validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 80.769



conv2d (Conv2D)	(None, 32, 32, 64)	1792	conv2d_6 (Conv2D)	(None, 8, 8, 256)	295168
conv2d_1 (Conv2D)	(None, 32, 32, 64)	36928	conv2d_7 (Conv2D)	(None, 8, 8, 256)	590080
conv2d_2 (Conv2D)	(None, 32, 32, 64)	36928	conv2d_8 (Conv2D)	(None, 8, 8, 256)	590080
max_pooling2d (MaxPooling2D)	(None, 16, 16, 64)	0	max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 256)	0
dropout (Dropout)	(None, 16, 16, 64)	0	dropout_2 (Dropout)	(None, 4, 4, 256)	0
conv2d_3 (Conv2D)	(None, 16, 16, 128)	73856	flatten (Flatten)	(None, 4096)	0
conv2d_4 (Conv2D)	(None, 16, 16, 128)	147584	dense (Dense)	(None, 1024)	4195328
conv2d_5 (Conv2D)	(None, 16, 16, 128)	147584	dropout_3 (Dropout)	(None, 1024)	0
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 128)	0	dense_1 (Dense)	(None, 10)	10250
<hr/>					
dropout_1 (Dropout)	(None, 8, 8, 128)	0	Total params:	6125578	(23.37 MB)
			Trainable params:	6125578	(23.37 MB)

Conclusiones y observaciones:

Aumentamos los filtros para ver como afectan al modelo. El accuracy apenas ha mejorado, pero hay más overfitting y el tiempo de entrenamiento a aumentado.

Aunque es cierto que no mucho(de 0:04:24 a 0:06:49), no considero que una mejora tan pequeña mereza la pena.

Proyecto VII

```

model = Sequential()
model.add(Conv2D(32, (3, 3), strides=1, padding='same', input_shape=(32, 32, 3))
model.add(Conv2D(32, (3, 3), strides=1, padding='same'))
model.add(Conv2D(32, (3, 3), strides=1, padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

model.add(Conv2D(64, (3, 3), strides=1, padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3), strides=1, padding='same'))
model.add(Conv2D(64, (3, 3), strides=1, padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))

```

conv2d_27 (Conv2D)	(None, 32, 32, 32)	896	conv2d_33 (Conv2D)	(None, 8, 8, 128)
conv2d_28 (Conv2D)	(None, 32, 32, 32)	9248	batch_normalization_10 (BatchNormalization)	(None, 8, 8, 128)
conv2d_29 (Conv2D)	(None, 32, 32, 32)	9248	conv2d_34 (Conv2D)	(None, 8, 8, 128)
max_pooling2d_9 (MaxPooling2D)	(None, 16, 16, 32)	0	conv2d_35 (Conv2D)	(None, 8, 8, 128)
dropout_12 (Dropout)	(None, 16, 16, 32)	0	max_pooling2d_11 (MaxPooling2D)	(None, 4, 4, 128)
conv2d_30 (Conv2D)	(None, 16, 16, 64)	18496	dropout_14 (Dropout)	(None, 4, 4, 128)
batch_normalization_9 (BatchNormalization)	(None, 16, 16, 64)	256	flatten_3 (Flatten)	(None, 2048)
conv2d_31 (Conv2D)	(None, 16, 16, 64)	36928	dense_6 (Dense)	(None, 500)
conv2d_32 (Conv2D)	(None, 16, 16, 64)	36928	dropout_15 (Dropout)	(None, 500)
max_pooling2d_10 (MaxPooling2D)	(None, 8, 8, 64)	0	dense_7 (Dense)	(None, 10)
dropout_13 (Dropout)	(None, 8, 8, 64)	0		5010
<hr/>				
			Total params:	1511046 (5.76 MB)
			Trainable params:	1510662 (5.76 MB)

```

history = model.fit(x_train_scaled, y_train,
                     callbacks=[modelchk, callback_loss, callback_acc],
                     epochs=300, batch_size= 512,
                     validation_data=(x_val_scaled, y_val))

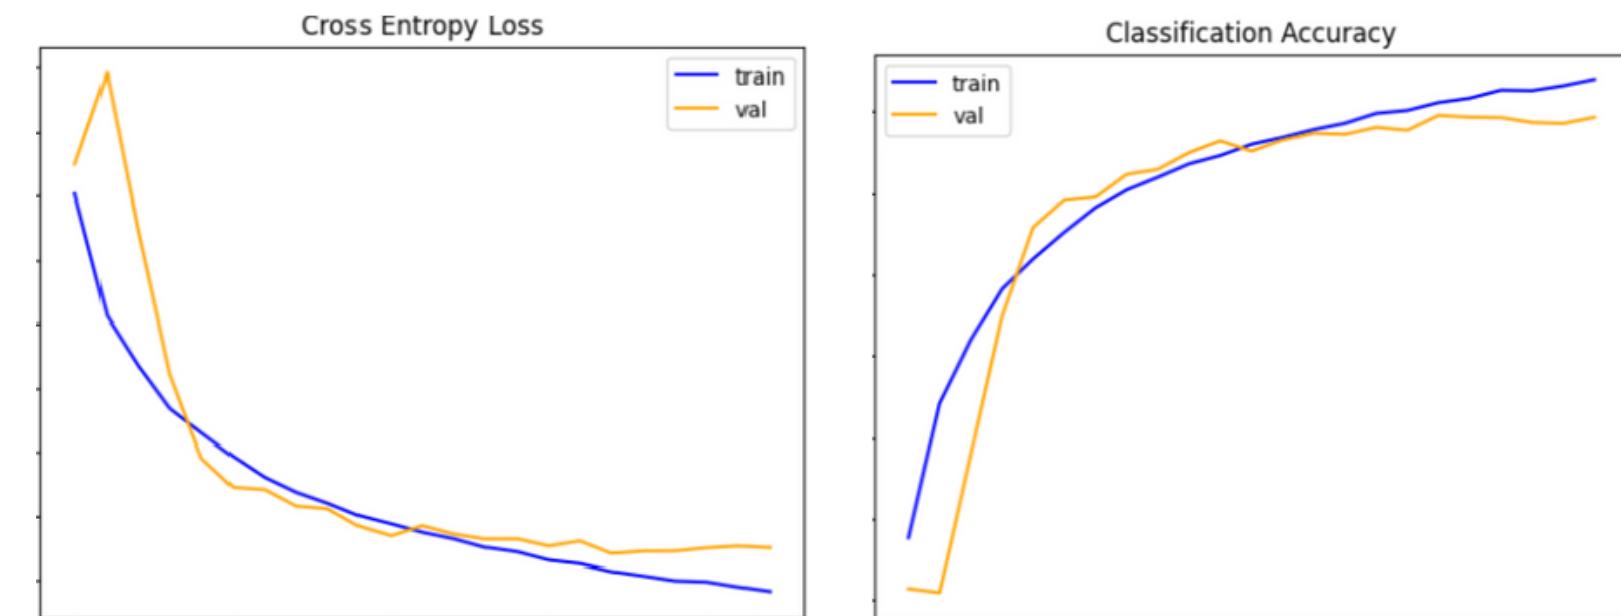
```

```

_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))

> 78.440

```



Conclusiones y observaciones:

Le añado 2 batchnormalization ya que no tenemos mucho overfitting y no queremos dificultar demasiado el aprendizaje del modelo.

Ha empeorado el modelo, probaremos con quitando uno ya que parece que el modelo aprende peor con 2.

Proyecto VIII

```
model = Sequential()
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same', input_shape=(32,32,3))
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

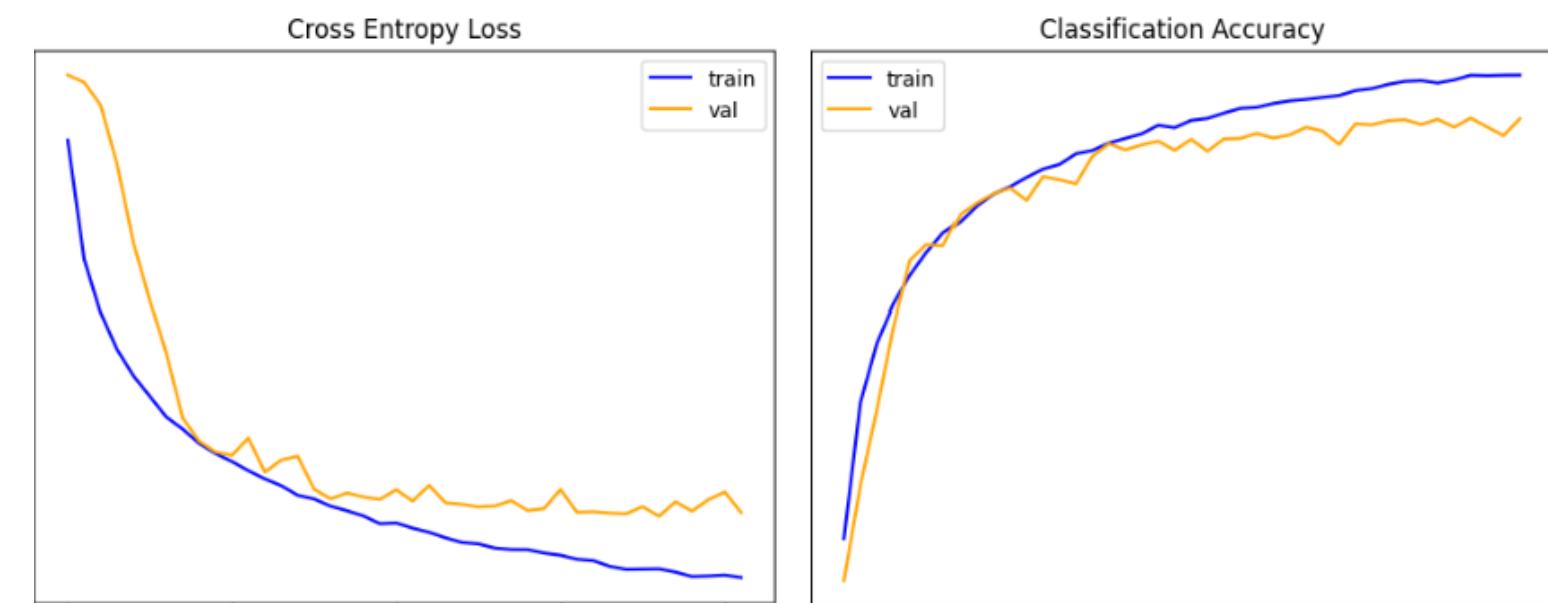
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))
```

```
history = model.fit(x_train_scaled, y_train,
                     callbacks=[modelchk, callback_loss, callback_acc],
                     epochs=300, batch_size= 512,
                     validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 81.600



Conclusiones y observaciones:

Le añado solamente un batchnormalization para no dificultar demasiado el aprendizaje.

Hasta ahora es el mejor resultado, su aprendizaje no se ha visto afectado en este caso.

=====			
conv2d_18 (Conv2D)	(None, 32, 32, 32)	896	batch_normalization_3 (BatchNormalization)
conv2d_19 (Conv2D)	(None, 32, 32, 32)	9248	conv2d_25 (Conv2D)
conv2d_20 (Conv2D)	(None, 32, 32, 32)	9248	conv2d_26 (Conv2D)
max_pooling2d_6 (MaxPooling2D)	(None, 16, 16, 32)	0	max_pooling2d_8 (MaxPooling2D)
dropout_8 (Dropout)	(None, 16, 16, 32)	0	dropout_10 (Dropout)
conv2d_21 (Conv2D)	(None, 16, 16, 64)	18496	flatten_2 (Flatten)
conv2d_22 (Conv2D)	(None, 16, 16, 64)	36928	dense_4 (Dense)
conv2d_23 (Conv2D)	(None, 16, 16, 64)	36928	dropout_11 (Dropout)
max_pooling2d_7 (MaxPooling2D)	(None, 8, 8, 64)	0	dense_5 (Dense)
dropout_9 (Dropout)	(None, 8, 8, 64)	0	=====
conv2d_24 (Conv2D)	(None, 8, 8, 128)	73856	Total params: 1510790 (5.76 MB)
			Trainable params: 1510534 (5.76 MB)

Proyecto IX

```
model = Sequential()
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same', input_shape=(32, 32, 3)))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(256, (3, 3), strides=1, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(256, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(256, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.25))

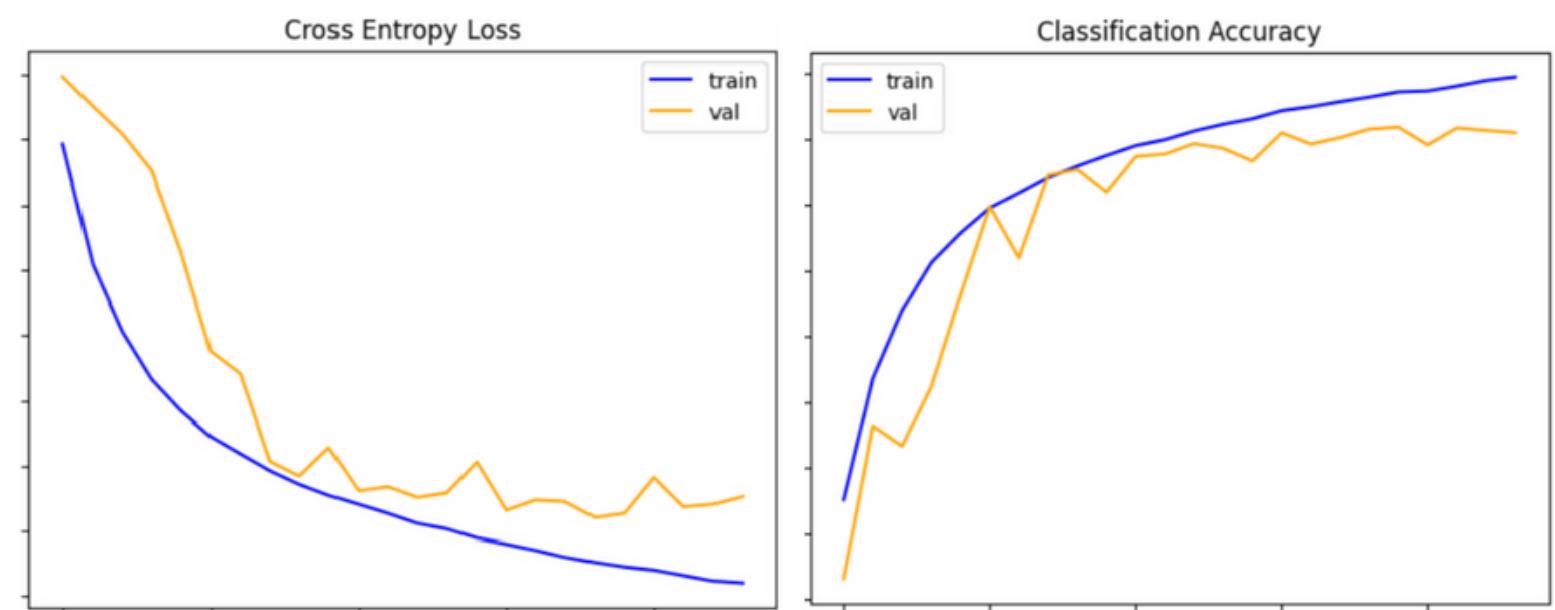
model.add(Flatten())
model.add(Dense(1024, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))
```

```
=====
conv2d_18 (Conv2D)    (None, 32, 32, 64)    1792      batch_normalization_3 (BatchNormalization)   1024
conv2d_19 (Conv2D)    (None, 32, 32, 64)    36928     conv2d_25 (Conv2D)    (None, 8, 8, 256)    590080
conv2d_20 (Conv2D)    (None, 32, 32, 64)    36928     conv2d_26 (Conv2D)    (None, 8, 8, 256)    590080
max_pooling2d_6 (MaxPooling2D) (None, 16, 16, 64) 0          max_pooling2d_8 (MaxPooling2D) (None, 4, 4, 256) 0
dropout_8 (Dropout)   (None, 16, 16, 64)    0          dropout_10 (Dropout)  (None, 4, 4, 256)    0
conv2d_21 (Conv2D)    (None, 16, 16, 128)   73856     flatten_2 (Flatten)  (None, 4096)       0
conv2d_22 (Conv2D)    (None, 16, 16, 128)   147584    dense_4 (Dense)   (None, 1024)      4195328
conv2d_23 (Conv2D)    (None, 16, 16, 128)   147584    dropout_11 (Dropout) (None, 1024)      0
max_pooling2d_7 (MaxPooling2D) (None, 8, 8, 128) 0          dense_5 (Dense)   (None, 10)        10250
dropout_9 (Dropout)   (None, 8, 8, 128)    0          =====
Total params: 6126602 (23.37 MB)
Trainable params: 6126090 (23.37 MB)
```

```
history = model.fit(x_train_scaled, y_train,
                     callbacks=[modelchk, callback_loss, callback_acc],
                     epochs=300, batch_size= 512,
                     validation_data=(x_val_scaled, y_val))
```

```
_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 80.700



Conclusiones y observaciones:

Por último, antes de pasar a técnicas mas avanzadas, probaremos como afecta el batchnormalization al notebook con aumento de filtros(proyecto VIII), ya que en ese caso necesitabamos mejorar el overfitting.

Ni el accuracy ni el overfitting mejoran, de momento el proyecto ganador es el X.

Proyecto X - DG

```

model = Sequential()
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same', input_shape=(32,32,3))
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))

```

```

=====
conv2d_7 (Conv2D)      (None, 32, 32, 32)    896      batch_normalization (Batch Normalization)      (None, 8, 8, 128)    512
conv2d_8 (Conv2D)      (None, 32, 32, 32)    9248     conv2d_14 (Conv2D)      (None, 8, 8, 128)    147584
conv2d_9 (Conv2D)      (None, 32, 32, 32)    9248     conv2d_15 (Conv2D)      (None, 8, 8, 128)    147584
max_pooling2d_2 (MaxPooling2D) (None, 16, 16, 32) 0        max_pooling2d_4 (MaxPooling2D) (None, 4, 4, 128) 0
dropout_2 (Dropout)   (None, 16, 16, 32)    0        dropout_4 (Dropout)   (None, 4, 4, 128)    0
conv2d_10 (Conv2D)    (None, 16, 16, 64)   18496     flatten (Flatten)   (None, 2048)      0
conv2d_11 (Conv2D)    (None, 16, 16, 64)   36928     dense (Dense)      (None, 500)      1024500
conv2d_12 (Conv2D)    (None, 16, 16, 64)   36928     dropout_5 (Dropout) (None, 500)      0
max_pooling2d_3 (MaxPooling2D) (None, 8, 8, 64) 0        dense_1 (Dense)    (None, 10)       5010
dropout_3 (Dropout)   (None, 8, 8, 64)    0        =====
conv2d_13 (Conv2D)    (None, 8, 8, 128)   73856     Total params: 1510790 (5.76 MB)
                                                               Trainable params: 1510534 (5.76 MB)

```

```

train_datagen = ImageDataGenerator(
    rescale=1./255, rotation_range = 15, horizontal_flip = True,
    width_shift_range=0.12, height_shift_range=0.12)

train_generator = train_datagen.flow(
    x_train, y_train, batch_size=30)

validation_datagen = ImageDataGenerator(
    rescale=1./255)
validation_generator = validation_datagen.flow(
    x_val, y_val, batch_size=20)

```

```

history = model.fit(train_generator, epochs=300, validation_data = validation_generator,
                     batch_size = 256, callbacks=[modelchk, callback_loss, callback_acc])

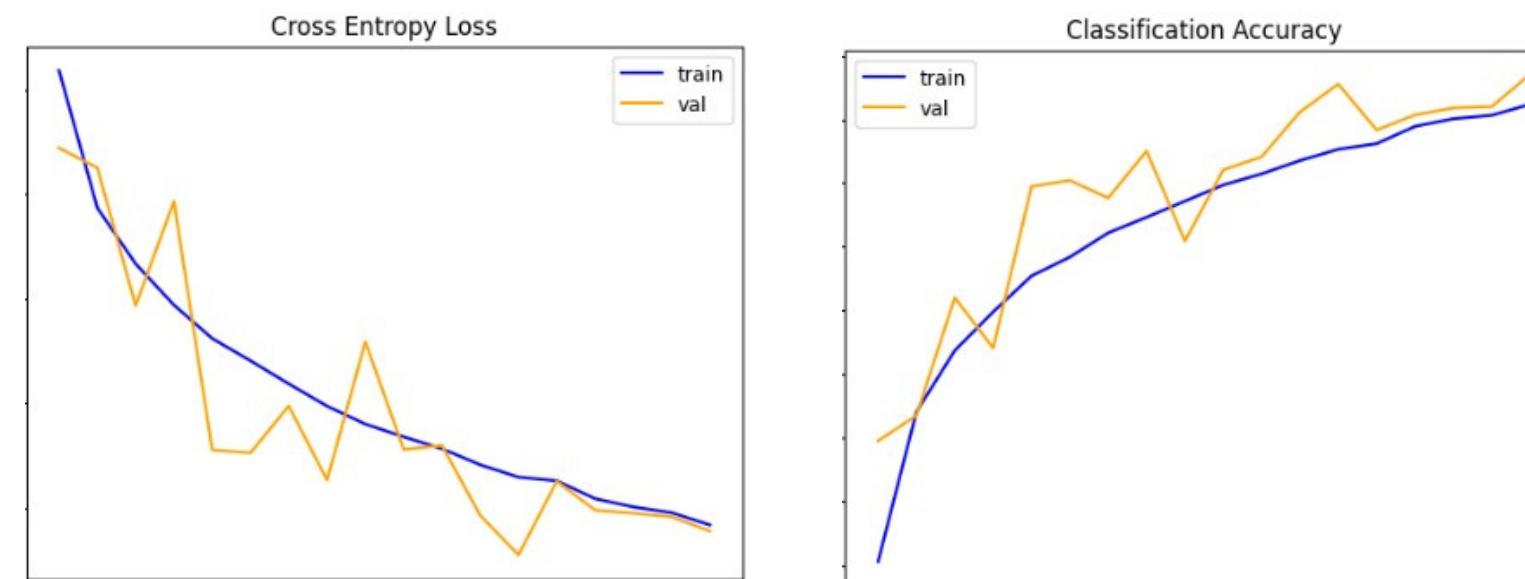
```

```

_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))

```

> 68.750



Conclusiones y observaciones:

Empezamos utilizando la técnica de data augmentation. En este primer intento el accuracy empeora drásticamente.

Vamos a intentar controlar el overfitting, parece que este modelo no funciona muy bien con los nuevos datos

Proyecto XI - DG

```

model = Sequential()
# Primera capa de convolución
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same', input_shape=(32,32,3))
model.add(BatchNormalization())
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(32, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))
# Segunda capa de convolución
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(64, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))
# Tercera capa de convolución
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(Conv2D(128, (3, 3), strides=1, activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))

```

```

=====
conv2d (Conv2D)          (None, 32, 32, 32)    896      conv2d_4 (Conv2D)          (None, 8, 8, 128)    73856
batch_normalization (Batch Normalization) (None, 32, 32, 32)    128      batch_normalization_2 (Batch Normalization) (None, 8, 8, 128)    512
conv2d_1 (Conv2D)          (None, 32, 32, 32)    9248     conv2d_7 (Conv2D)          (None, 8, 8, 128)    147584
conv2d_2 (Conv2D)          (None, 32, 32, 32)    9248     conv2d_8 (Conv2D)          (None, 8, 8, 128)    147584
max_pooling2d (MaxPooling2D) (None, 16, 16, 32)    0        max_pooling2d_2 (MaxPooling2D) (None, 4, 4, 128)    0
dropout (Dropout)          (None, 16, 16, 32)    0        dropout_2 (Dropout)        (None, 4, 4, 128)    0
conv2d_3 (Conv2D)          (None, 16, 16, 64)    18496     flatten (Flatten)        (None, 2048)      0
batch_normalization_1 (Batch Normalization) (None, 16, 16, 64)    256      dense (Dense)          (None, 500)      1024500
conv2d_4 (Conv2D)          (None, 16, 16, 64)    36928     dropout_3 (Dropout)        (None, 500)      0
conv2d_5 (Conv2D)          (None, 16, 16, 64)    36928     dense_1 (Dense)          (None, 10)       5010
max_pooling2d_1 (MaxPooling2D) (None, 8, 8, 64)    0        =====
Total params: 1511174 (5.76 MB)
Trainable params: 1510726 (5.76 MB)
Non-trainable params: 448 (1.75 KB)
dropout_1 (Dropout)        (None, 8, 8, 64)    0

```

```

train_datagen = ImageDataGenerator(
    rescale=1./255, rotation_range = 15, horizontal_flip = True,
    width_shift_range=0.12, height_shift_range=0.12)

train_generator = train_datagen.flow(
    x_train, y_train, batch_size=30)

```

```

validation_datagen = ImageDataGenerator(
    rescale=1./255)
validation_generator = validation_datagen.flow(
    x_val, y_val, batch_size=20)

```

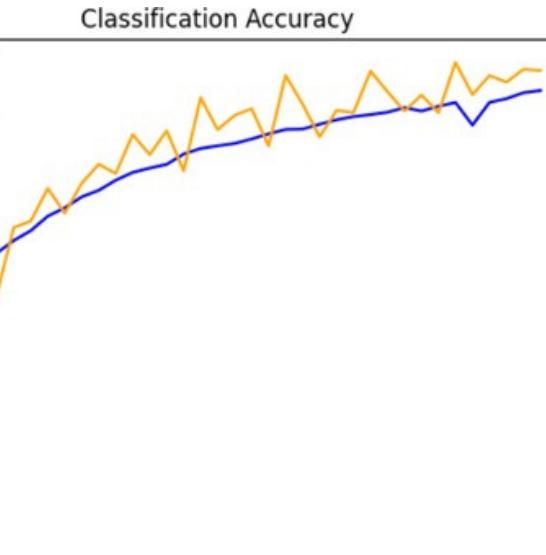
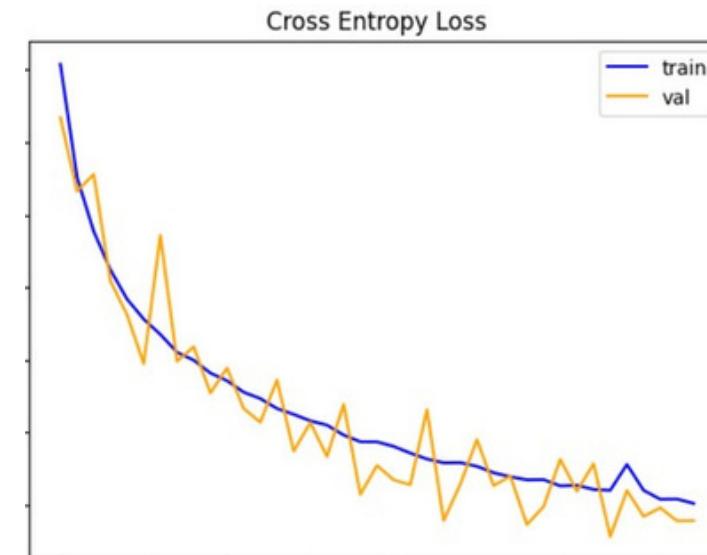
```

history = model.fit(train_generator, epochs=300, validation_data = validation_generator,
                     batch_size = 256, callbacks=[modelchk, callback_loss, callback_acc])

_, acc = model.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))

```

> 82.250



Conclusiones y observaciones:

Añadimos batchnormalization. El resultado es un accuracy de 82, el mejor modelo hasta ahora.

Aunque el segundo mejor modelo era de 81, y tardaba 4min aproximadamente en cargar, este en cambio ha tardado 24min por lo que no merece la pena una mejora tan pequeña.

Proyecto XII - TL

```
model_vgg16 = vgg16.VGG16(include_top = False, weights = 'imagenet',
                           input_shape = (32,32,3))

model_post_vgg = ks.Sequential()

model_post_vgg.add(final_vgg16)
model_post_vgg.add(ks.layers.Dense(500, activation='relu'))
model_post_vgg.add(ks.layers.Dropout(0.5))
model_post_vgg.add(ks.layers.Dense(10, activation='softmax'))

model_post_vgg.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
model (Functional)	(None, 512)	14714688
dense (Dense)	(None, 500)	256500
dropout (Dropout)	(None, 500)	0
dense_1 (Dense)	(None, 10)	5010
<hr/>		

Total params: 14976198 (57.13 MB)
Trainable params: 14716038 (56.14 MB)
Non-trainable params: 260160 (1016.25 KB)

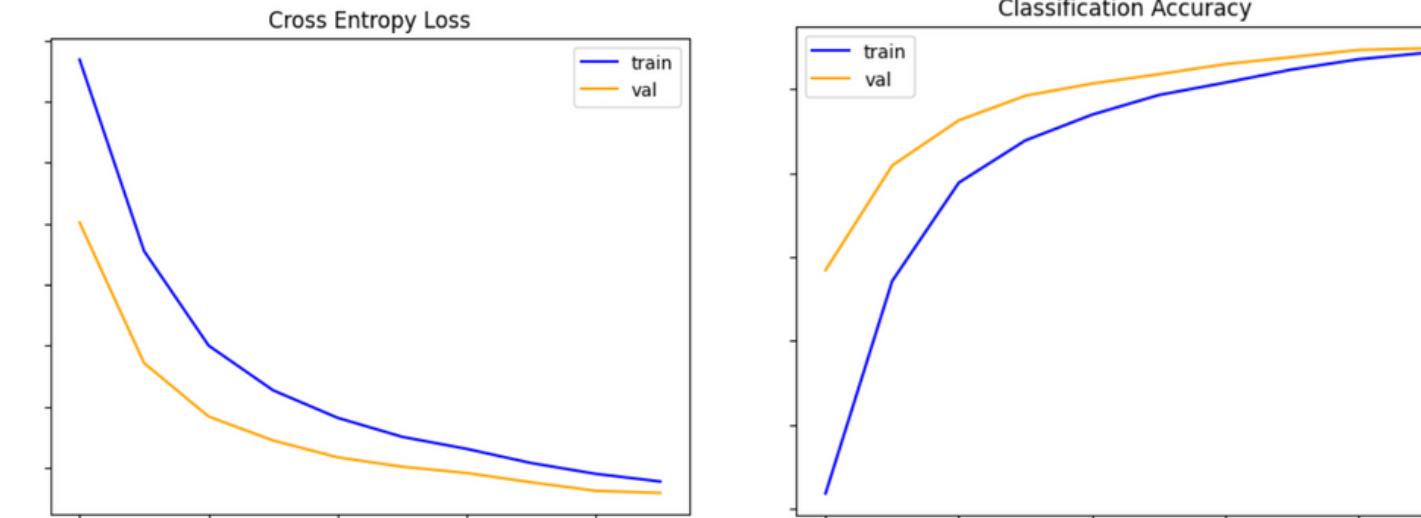
```
new_adam = Adam(learning_rate=0.000005)
model_post_vgg.compile(loss='sparse_categorical_crossentropy',
                        optimizer=new_adam, metrics=['accuracy'])
```

```
history = model_post_vgg.fit(x=x_train_scaled, y=y_train, batch_size=512,
                               epochs=10, callbacks=[callback],
                               validation_data=(x_val_scaled, y_val))
```

Tiempo de entrenamiento: 0:00:00.007997

```
, acc = model_post_vgg.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 74.190



Conclusiones y observaciones:

He implementado transfer learning con el modelo vgg16, además he cambiado el learning rate.

El accuracy ha empeorado, pero el tiempo de entrenamiento ha disminuido muchísimo, por lo que voy a intentar mejorar el modelo a partir de aquí.

Proyecto XIII - TL + WD

```
model_vgg16 = vgg16.VGG16(include_top = False, weights = 'imagenet',
                           input_shape = (32,32,3))
```

```
model_post_vgg = ks.Sequential()

model_post_vgg.add(final_vgg16)
model_post_vgg.add(Dense(500, activation='relu', kernel_regularizer=l2(0.00001)))
model_post_vgg.add(Dropout(0.5))
model_post_vgg.add(Dense(10, activation='softmax'))
```

Model: "sequential_1"

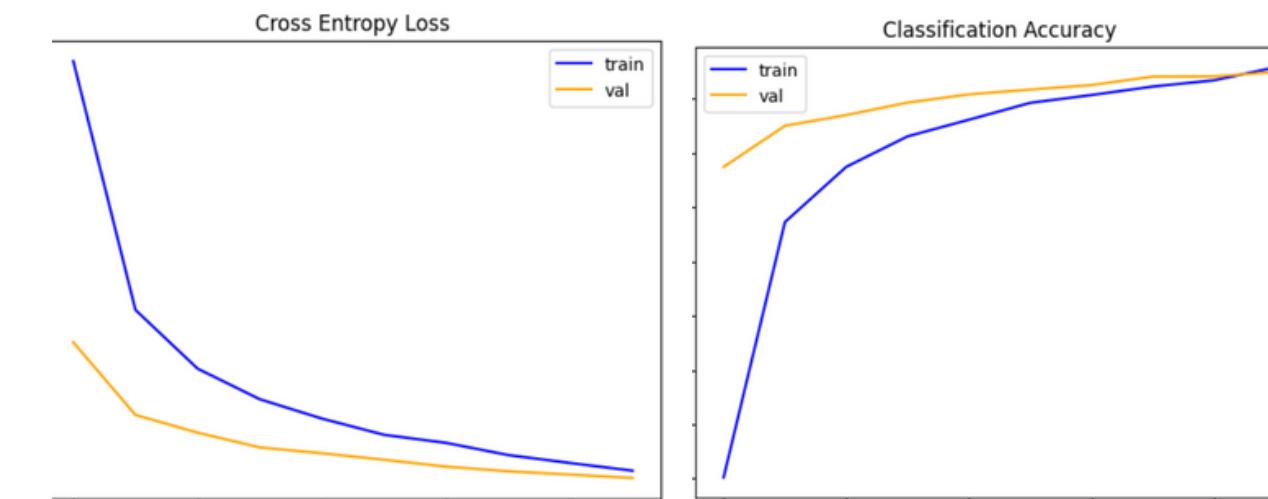
Layer (type)	Output Shape	Param #
<hr/>		
model (Functional)	(None, 512)	14714688
dense_2 (Dense)	(None, 500)	256500
dropout_1 (Dropout)	(None, 500)	0
dense_3 (Dense)	(None, 10)	5010
<hr/>		
Total params: 14976198 (57.13 MB)		
Trainable params: 14716038 (56.14 MB)		
Non-trainable params: 260160 (1016.25 KB)		

```
new_adam = Adam(learning_rate=0.000005)
model_post_vgg.compile(loss='sparse_categorical_crossentropy',
                       optimizer=new_adam, metrics=['accuracy'])
```

```
history = model_post_vgg.fit(x=x_train_scaled, y=y_train, batch_size=512,
                             epochs=10, callbacks=[callback],
                             validation_data=(x_val_scaled, y_val))
```

```
_ , acc = model_post_vgg.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 77.500



Conclusiones y observaciones:

He añadido el l2 con la intención de controlar el overfitting

Ahora toca cambiar el optimizador ya que Adam no funciona bien con esta técnica.

Proyecto XIV - TL + WD

```
model_vgg16 = vgg16.VGG16(include_top = False, weights = 'imagenet',
                           input_shape = (32,32,3))
```

```
model_post_vgg = ks.Sequential()

model_post_vgg.add(final_vgg16)
model_post_vgg.add(Dense(500, activation='relu', kernel_regularizer=l2(0.00001)))
model_post_vgg.add(Dropout(0.5))
model_post_vgg.add(Dense(10, activation='softmax'))
```

Model: "sequential_1"

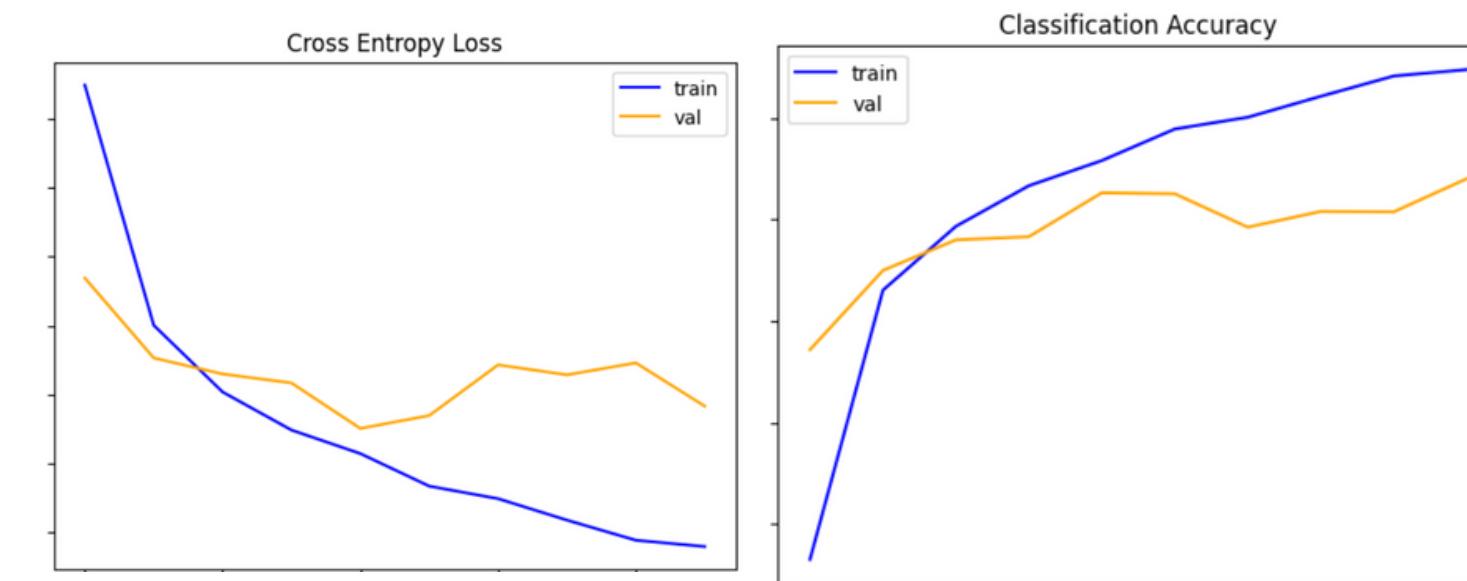
Layer (type)	Output Shape	Param #
=====		
model (Functional)	(None, 512)	14714688
dense_2 (Dense)	(None, 500)	256500
dropout_1 (Dropout)	(None, 500)	0
dense_3 (Dense)	(None, 10)	5010
=====		
Total params:	14976198 (57.13 MB)	
Trainable params:	14716038 (56.14 MB)	
Non-trainable params:	260160 (1016.25 KB)	

```
new_sgd = SGD(learning_rate=0.01, momentum=0.9)
model_post_vgg.compile(loss='sparse_categorical_crossentropy',
                       optimizer=new_sgd, metrics=['accuracy'])
```

```
history = model_post_vgg.fit(x=x_train_scaled, y=y_train, batch_size=512,
                             epochs=10, callbacks=[callback],
                             validation_data=(x_val_scaled, y_val))
```

```
_, acc = model_post_vgg.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 83.380



Conclusiones y observaciones:

Cambiamos el optimizador a SGD lo cual mejora el modelo.

Ahora queremos controlar el overfitting y para ello probaré con data generation, ya que parece que el modelo es muy complejo y tenemos pocos datos.

Proyecto XV - TL + DG

```
model_vgg16 = vgg16.VGG16(include_top = False, weights = 'imagenet',
                           input_shape = (32,32,3))
```

```
model_post_vgg = ks.Sequential()
```

```
model_post_vgg.add(final_vgg16)
model_post_vgg.add(Dense(500, activation='relu', kernel_regularizer=l2(0.00001)))
model_post_vgg.add(Dropout(0.5))
model_post_vgg.add(Dense(10, activation='softmax'))
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
<hr/>		
model (Functional)	(None, 512)	14714688
dense_2 (Dense)	(None, 500)	256500
dropout_1 (Dropout)	(None, 500)	0
dense_3 (Dense)	(None, 10)	5010
<hr/>		

Total params: 14976198 (57.13 MB)
Trainable params: 14716038 (56.14 MB)
Non-trainable params: 260160 (1016.25 KB)

```
new_sgd = SGD(learning_rate=0.01, momentum=0.9)
model_post_vgg.compile(loss='sparse_categorical_crossentropy',
                       optimizer=new_sgd, metrics=['accuracy'])
```

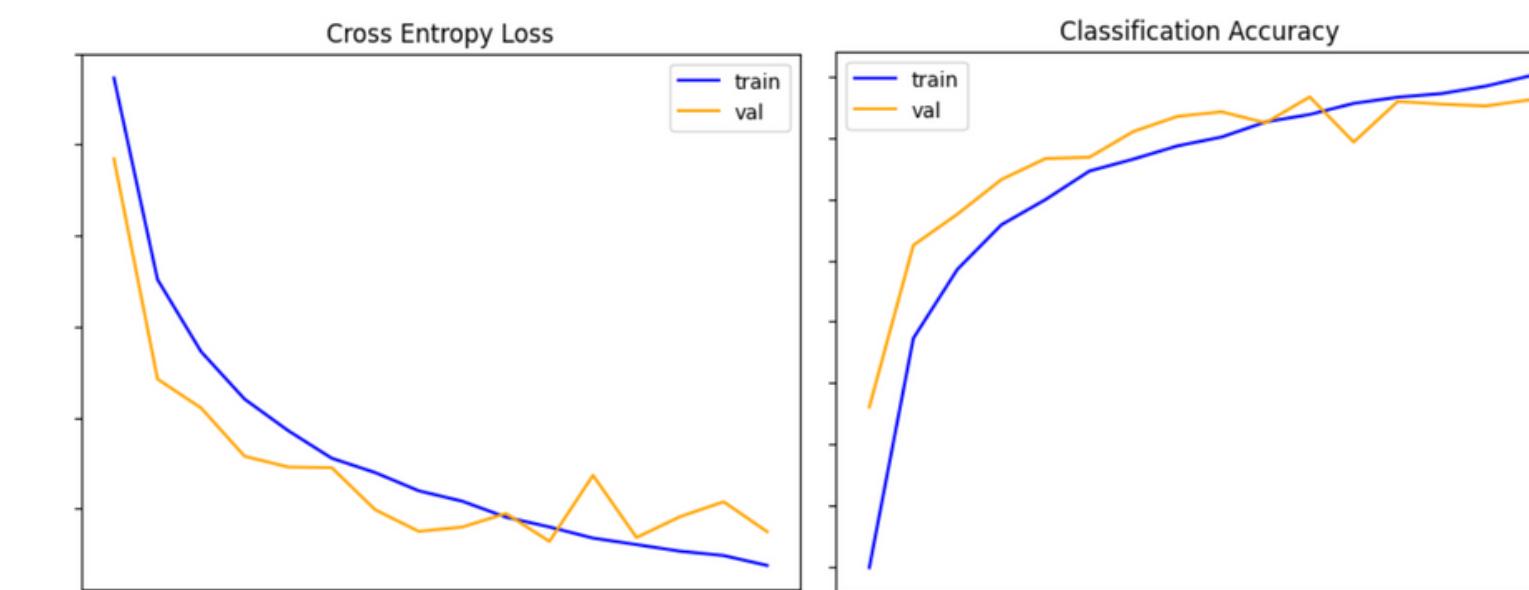
```
train_datagen = ImageDataGenerator(rescale=1./255,
                                    rotation_range = 15, horizontal_flip = True,
                                    width_shift_range=0.12,height_shift_range=0.12)
```

```
train_generator = train_datagen.flow(
    x_train, y_train,batch_size=30)
```

```
history = model_post_vgg.fit(train_generator, epochs=300,
                             validation_data = validation_generator,
                             batch_size = 256, callbacks=[callback])
```

```
_, acc = model_post_vgg.evaluate(x_test_scaled, y_test, verbose=0)
print('> %.3f' % (acc * 100.0))
```

> 82.720



Conclusiones y observaciones:

He aumentado también los epochs para sacar el máximo partido al modelos.

Parece que la teoría se ha confirmado, aumentar los datos ha controlado el overfitting.

Conclusión

Proyecto	Procedimiento	Accuracy	Train time
P.1	Aumento de neuronas	65,86	0:00:43
P.2	Aumento de capas	73,96	0:00:48
P.3	Aumento de convulciones	76,47	0:02:19
P.4	Dropouts 0.4	73,48	0:04:03
P.5	Dropouts 0.3	80,38	0:04:24
P.6	Cambiar los filtros	80,76	0:06:49
P.7	BatchNormalization v.1	78,44	0:02:59
P.8	BatchNormalization v.2	81,60	0:04:57
P.9	Filtros + BN	80,70	0:05:59
P.10	Datageneration	68,75	0:11:16
P.11	Datagen. + control overfitting	82,25	0:24:29
P.12	Transfer learning	74,19	0:00:00
P.13	T. learning + Weight Decay	77,50	0:00:00
P.14	Cambio optimizador	83,38	0:00:00
P.15	T. learning + D. Generation	82,72	0:00:00

Conclusión

Aunque hemos tenido varios modelos con un buen accuracy, nos hemos encontrado con 2 problemas principales, el overfitting y el tiempo de entrenamiento.

Al llegar a transfer learnign he podido apreciar lo eficaz que es para disminuir el segundo problema (el tiempo) por lo que he decidido partir de ahí mejorando el modelo, y al final el overfitting.

Sin duda podemos concluir que el proyecto 15 con un acc de 82,72 (el segundo más alto) 0 seg de tiempo y un overfitting controlado es el ganador (por encima del 14 que aunque tenía mejor accuracy también tenía problemas de sobreajuste).